

Block B: Computer Vision & Robotics


Monika Stangenberg 231648

Problem Overview

Client: Netherlands Plant Eco-phenotyping Centre

Automatic segmentation of plant roots from images provided by NPEC

Precisely controlled robot to perform in-vitro inoculation of plants with microbes or synthetic communities



Proposed AI-Powered Solution




Implementing CV pipeline which combines image segmentation techniques and deep learning




Training a model which identifies plants' roots from images provided by the client.



Developing and benchmarking PID and RL controllers to precisely target plant roots using robotics pipettes.



Integrating computer vision outputs to map root tips' (pixel coordinates) to robotic arm positions.




Benefits for NPEC

- Help with identifying and tracking root growth.
- Automating process of measuring plants' roots, which is originally very time consuming.
- Precise and automated root targeting by robot (Hades).

Pre-processing

- Trimming 150 pixels from right side, because images have white strip there
- Using middle x value to find valid columns (based on pixel values)
- Using middle y value to find valid rows (based on pixel values)
- Checking if the image is square, if not cropping excess rows/columns to make the image square




```
def trim_picture(picture, mask):
    # Trim 150 pixels from the right side of the image
    picture = picture[:, :-150, :]
    mask = mask[:, :-150]
    # Original dimensions of the picture
    xo, yo, zo = picture.shape
    # I find the middle row and identify valid columns based on pixel values
    x1 = xo//2
    valid_columns = (picture[x1, :, :] >= 100).any(axis=1)
    trimmed_picture1 = picture[:, valid_columns, :] # Trim invalid columns
    trimmed_mask1 = mask[:, valid_columns] # Trim invalid columns
    # New dimensions of the picture
    x_trimmed, y_trimmed, z_trimmed = trimmed_picture1.shape
    # I find the middle row and identify valid columns based on pixel values
    y2 = y_trimmed//2
    valid_rows = (trimmed_picture1[:, y2, :] >= 100).any(axis=1)
    trimmed_picture2 = trimmed_picture1[valid_rows, :, :]
    trimmed_mask2 = trimmed_mask1[valid_rows, :]
    # New dimensions of the picture
    x, y, z = trimmed_picture2.shape
    # Checking if the picture is already a square
    if x == y:
        return trimmed_picture2, trimmed_mask2
    min_dim = min(x, y)
    if x > y: # Crop excess rows if height is greater than width
        excess = x - min_dim
        start = excess // 2
        end = start + min_dim
        squared_picture = trimmed_picture2[start:end, :, :]
        squared_mask = trimmed_mask2[start:end, :]
    else: # Crop excess columns if width is greater than height
        excess = y - min_dim
        start = excess // 2
        end = start + min_dim
        squared_picture = trimmed_picture2[:, start:end, :]
        squared_mask = trimmed_mask2[:, start:end]
    return squared_picture, squared_mask
```


Pre-processing Patchifying

- patch size = 256
- I used padding to make sure that dimensions of all images and masks are divisible by patch size
- I save patched images with unique filename (001, 002, ...)

Cropped Image with Grid



Model 1 Results

- Best score:
 - Accuracy/val accuracy: 0.9896/0.9878
 - Loss/val loss: 0.0545/0.0704
 - I did not include f1 score

Because of the lack of a GPU, my computer used the CPU to train the model, which significantly increased the training time.

```
Epoch 1/10
1310/1310 [=====] - 9944s 8s/step - loss: 0.1217 - accuracy: 0.9791 - val_loss: 0.0758 - val_accuracy: 0.9860
Epoch 2/10
1310/1310 [=====] - 10236s 8s/step - loss: 0.0586 - accuracy: 0.9886 - val_loss: 0.0670 - val_accuracy: 0.9869
Epoch 3/10
1310/1310 [=====] - 9355s 7s/step - loss: 0.0567 - accuracy: 0.9890 - val_loss: 0.1037 - val_accuracy: 0.9864
Epoch 4/10
1310/1310 [=====] - 9458s 7s/step - loss: 0.0555 - accuracy: 0.9892 - val_loss: 0.1681 - val_accuracy: 0.9861
Epoch 5/10
1310/1310 [=====] - 9457s 7s/step - loss: 0.0558 - accuracy: 0.9892 - val_loss: 0.0677 - val_accuracy: 0.9870
Epoch 6/10
1310/1310 [=====] - 9030s 7s/step - loss: 0.0556 - accuracy: 0.9893 - val_loss: 0.0626 - val_accuracy: 0.9879
Epoch 7/10
1310/1310 [=====] - 9120s 7s/step - loss: 0.0550 - accuracy: 0.9894 - val_loss: 0.0636 - val_accuracy: 0.9877
Epoch 8/10
1310/1310 [=====] - 9087s 7s/step - loss: 0.0548 - accuracy: 0.9894 - val_loss: 0.1054 - val_accuracy: 0.9806
Epoch 9/10
1310/1310 [=====] - 49356s 38s/step - loss: 0.0545 - accuracy: 0.9895 - val_loss: 0.0920 - val_accuracy: 0.9873
Epoch 10/10
1310/1310 [=====] - 36887s 28s/step - loss: 0.0545 - accuracy: 0.9896 - val_loss: 0.0704 - val_accuracy: 0.9878
```

```
699m 4.4s
Epoch 1/10
1310/1310 [=====]
```

Model 1 Architecture

```
model = Sequential()

# input layer: dense layer helps reduce dimensions
model.add(Dense(200, input_shape=(256, 256, 1), activation='relu')) # input shape (256, 256, 1)

# downsampling 1
model.add(Conv2D(32, (3, 3), activation='relu', padding='same'))
model.add(BatchNormalization())
model.add(MaxPooling2D((2, 2)))
model.add(Dropout(0.3))

# downsampling 2
model.add(Conv2D(64, (3, 3), activation='relu', padding='same'))
model.add(BatchNormalization())
model.add(MaxPooling2D((2, 2)))
model.add(Dropout(0.3))

# bottleneck
model.add(Conv2D(128, (3, 3), activation='relu', padding='same'))
model.add(BatchNormalization())

# upsampling 1
model.add(UpSampling2D((2, 2)))
model.add(Conv2D(64, (3, 3), activation='relu', padding='same'))
model.add(BatchNormalization())

# upsampling 2
model.add(UpSampling2D((2, 2)))
model.add(Conv2D(32, (3, 3), activation='relu', padding='same'))
model.add(BatchNormalization())

# output layer
model.add(Conv2D(1, (1, 1), activation='sigmoid'))
```

To meet client requirements,
I used "val_..." images for
validation and Y2B_24 data
for training.

```
model.compile(optimizer='adam',
              loss='binary_crossentropy', # for binary classification
              metrics=['accuracy'])
```

```
# defining early stopping
early_stopping = EarlyStopping(monitor='val_loss', patience=5, restore_best_weights=True)


# training the model
history = model.fit(
    x_train, y_train,
    validation_data=(x_test, y_test),
    epochs=10,          # only 10 epochs, because of long training time
    batch_size=16,
    callbacks=[early_stopping] # early stopping to prevent overfitting (but it is not real)
```

- Early stopping is not necessary, since I am training only with 10 epochs
- I am using a convolutional autoencoder

Model 1 Template Image Results

The model is detecting roots quite well. However, in some places they are not connected. It is also detecting background as root system part.

Inference template image →



Model 1

Visual Results

```
kernel = np.ones((5, 5), dtype="uint8")  
im_rod_d1 = cv2.dilate(predicted_mask, kernel, iterations=2)  
im_rod_closing = cv2.erode(im_rod_d1, kernel, iterations=1)
```

Original Image



Predicted Mask




Mask after closing

- Detecting big amount of background noise
- Parts of roots are not connected with each other
- I am using closing to remove small gaps in masks

Model 1

Visual Results



```
def generate_bounding_boxes(mask, min_area=230, y_range=(370, 730)):  
  
    contours, _ = cv2.findContours(mask, cv2.RETR_EXTERNAL, cv2.CHAIN_APPROX_SIMPLE)  
    bboxes = []  
    for contour in contours:  
        if cv2.contourArea(contour) >= min_area:  
            x, y, w, h = cv2.boundingRect(contour)  
            if y_range is None or (y_range[0] <= y <= y_range[1]):  
                bboxes.append((x, y, w, h))  
    return bboxes
```

Because of gaps in root masks, I can not detect roots correctly.

I detect objects on the masks and create bounding boxes around the 5 tallest ones.

I set a minimum size to avoid detecting background noise as roots in pictures where there are fewer than 5 roots.

Model 1 Bounding Boxes

Bounding boxes need to start in
 $y_range = (370, 730)$


I implemented this solution
because I noticed that in some
cases a bounding box with
background noise was detected as
one of the 5 largest boxes.

This appeared mainly in images
where the roots were very small.
By limiting the space where
bounding boxes could start, I was
able to target the roots more
effectively.


```
def generate_bounding_boxes(mask, min_area=230, y_range=(370, 730)):

    contours, _ = cv2.findContours(mask, cv2.RETR_EXTERNAL, cv2.CHAIN_APPROX_SIMPLE)
    bboxes = []
    for contour in contours:
        if cv2.contourArea(contour) >= min_area:
            x, y, w, h = cv2.boundingRect(contour)
            if y_range is None or (y_range[0] <= y <= y_range[1]):
                bboxes.append((x, y, w, h))
    return bboxes
```

Before implementation





Original image



Model 1

Bounding Boxes

- To process root images I skeletonize them, analyzed their structures and extracted metrics.
- I processes CSV files containing branch data extracted from test images. It calculates the maximum root length for each plant in every image from Kaggle submission photos and generates a consolidated CSV file.
- Kaggle score:

Root length prediction BUas & NPEC Y2B 2024-25							Late Submission		...
Overview	Data	Discussion	Leaderboard	Rules	Team	Submissions			
54	▲ 2	Nils Vos				25.257	7	1d	
55	▲ 16	Monika Stangenberg				25.583	2	1d	

Model 2 structure

I used model provided in self-study notebook 7

I also added f1 to metrics.

Because of long training time, epochs = 6.

```
model.compile(
    optimizer='adam',
    loss='binary_crossentropy',
    metrics=['accuracy', f1_score]
)
```

```
history = model.fit(
    x_train, y_train,
    validation_data=(x_test, y_test),
    epochs=6,
    batch_size=16,
    callbacks=[early_stopping]
)
```

```
model = Sequential()

model.add(Conv2D(16, (3, 3), activation='relu', padding='same', input_shape=(256, 256, 1)))
model.add(Dropout(0.1))
model.add(Conv2D(16, (3, 3), activation='relu', padding='same'))
model.add(MaxPooling2D((2, 2)))

model.add(Conv2D(32, (3, 3), activation='relu', padding='same'))
model.add(Dropout(0.1))
model.add(Conv2D(32, (3, 3), activation='relu', padding='same'))
model.add(MaxPooling2D((2, 2)))

model.add(Conv2D(64, (3, 3), activation='relu', padding='same'))
model.add(Dropout(0.2))
model.add(Conv2D(64, (3, 3), activation='relu', padding='same'))
model.add(MaxPooling2D((2, 2)))

model.add(Conv2D(128, (3, 3), activation='relu', padding='same'))
model.add(Dropout(0.2))
model.add(Conv2D(128, (3, 3), activation='relu', padding='same'))
model.add(MaxPooling2D((2, 2)))

model.add(Conv2D(256, (3, 3), activation='relu', padding='same'))
model.add(Dropout(0.3))
model.add(Conv2D(256, (3, 3), activation='relu', padding='same'))

model.add(Conv2DTranspose(128, (2, 2), strides=(2, 2), padding='same'))
model.add(Conv2D(128, (3, 3), activation='relu', padding='same'))
model.add(Dropout(0.2))
model.add(Conv2D(128, (3, 3), activation='relu', padding='same'))

model.add(Conv2DTranspose(64, (2, 2), strides=(2, 2), padding='same'))
model.add(Conv2D(64, (3, 3), activation='relu', padding='same'))
model.add(Dropout(0.2))
model.add(Conv2D(64, (3, 3), activation='relu', padding='same'))

model.add(Conv2DTranspose(32, (2, 2), strides=(2, 2), padding='same'))
model.add(Conv2D(32, (3, 3), activation='relu', padding='same'))
model.add(Dropout(0.1))
model.add(Conv2D(32, (3, 3), activation='relu', padding='same'))

model.add(Conv2DTranspose(16, (2, 2), strides=(2, 2), padding='same'))
model.add(Conv2D(16, (3, 3), activation='relu', padding='same'))
model.add(Dropout(0.1))
model.add(Conv2D(16, (3, 3), activation='relu', padding='same'))

model.add(Conv2D(1, (1, 1), activation='sigmoid'))
```


Model 2



Visual Results

This model didn't work well.
Output is just plain black screen.

The model structure is correct.

Loss function should be changed.

Because of this results, I didn't
continue working on this model.



```
- loss: 0.0797 - accuracy: 0.9868 - f1_score: 4.4571e-06 - val_loss: 0.0323 - val_accuracy: 0.9962 - val_f1_score: 0.0000e+00
- loss: 0.0696 - accuracy: 0.9874 - f1_score: 0.0000e+00 - val_loss: 0.0293 - val_accuracy: 0.9962 - val_f1_score: 0.0000e+00
- loss: 0.0695 - accuracy: 0.9874 - f1_score: 0.0000e+00 - val_loss: 0.0253 - val_accuracy: 0.9962 - val_f1_score: 0.0000e+00
- loss: 0.0692 - accuracy: 0.9874 - f1_score: 0.0000e+00 - val_loss: 0.0258 - val_accuracy: 0.9962 - val_f1_score: 0.0000e+00
- loss: 0.0691 - accuracy: 0.9874 - f1_score: 0.0000e+00 - val_loss: 0.0356 - val_accuracy: 0.9962 - val_f1_score: 0.0000e+00
- loss: 0.0688 - accuracy: 0.9874 - f1_score: 0.0000e+00 - val_loss: 0.0275 - val_accuracy: 0.9962 - val_f1_score: 0.0000e+00
```

Model 3

I didn't include f1 score in my previous working model.

Since my model training time is very long, I decided to make my dataset smaller to decrease training time.

```
model = Sequential()

# input layer: dense layer helps reduce dimensions
model.add(Dense(200, input_shape=(256, 256, 1), activation='relu')) # input shape (256, 256, 1)

# downsampling 1
model.add(Conv2D(32, (3, 3), activation='relu', padding='same'))
model.add(BatchNormalization())
model.add(MaxPooling2D((2, 2)))
model.add(Dropout(0.3))

# downsampling 2
model.add(Conv2D(64, (3, 3), activation='relu', padding='same'))
model.add(BatchNormalization())
model.add(MaxPooling2D((2, 2)))
model.add(Dropout(0.3))

# bottleneck
model.add(Conv2D(128, (3, 3), activation='relu', padding='same'))
model.add(BatchNormalization())

# upsampling 1
model.add(UpSampling2D((2, 2)))
model.add(Conv2D(64, (3, 3), activation='relu', padding='same'))
model.add(BatchNormalization())

# upsampling 2
model.add(UpSampling2D((2, 2)))
model.add(Conv2D(32, (3, 3), activation='relu', padding='same'))
model.add(BatchNormalization())

# output layer
model.add(Conv2D(1, (1, 1), activation='sigmoid'))
```

```
train_pairs = find_image_mask_pairs(train_image_dir, train_mask_dir)
x_train_paths = [pair[0] for pair in train_pairs]
y_train_paths = [pair[1] for pair in train_pairs]

test_pairs = find_image_mask_pairs(test_image_dir, test_mask_dir)
x_test_paths = [pair[0] for pair in test_pairs]
y_test_paths = [pair[1] for pair in test_pairs]

subset_size = 10000

x_train_paths = x_train_paths[:subset_size]
y_train_paths = y_train_paths[:subset_size]



# Load the datasets
x_train = load_images(x_train_paths)
y_train = load_images(y_train_paths)
x_test = load_images(x_test_paths)
y_test = load_images(y_test_paths)
```

Model architecture is the same as in model 1

Model 3 Results

```
best_val_loss = min(history.history['val_loss'])
best_val_f1 = max(history.history['val_f1'])
print(f"Best validation loss: {best_val_loss}")
print(f"Best validation f1: {best_val_f1}")
```

Best validation loss: 0.016697853803634644
Best validation f1: 0.6318885087966919




```
Epoch 1/7
625/625 [=====] - 5243s 8s/step - loss: 0.1802 - accuracy: 0.9745 - f1: 0.1050 - val_loss: 0.0242 - val_accuracy: 0.9967 - val_f1: 0.3364
Epoch 2/7
625/625 [=====] - 7354s 12s/step - loss: 0.0617 - accuracy: 0.9879 - f1: 0.3975 - val_loss: 0.0167 - val_accuracy: 0.9969 - val_f1: 0.2528
Epoch 3/7
625/625 [=====] - 5260s 8s/step - loss: 0.0599 - accuracy: 0.9883 - f1: 0.4826 - val_loss: 0.0174 - val_accuracy: 0.9974 - val_f1: 0.4571
Epoch 4/7
625/625 [=====] - 5331s 9s/step - loss: 0.0596 - accuracy: 0.9883 - f1: 0.4884 - val_loss: 0.0247 - val_accuracy: 0.9962 - val_f1: 0.0050
Epoch 5/7
625/625 [=====] - 4664s 7s/step - loss: 0.0591 - accuracy: 0.9885 - f1: 0.5181 - val_loss: 0.0353 - val_accuracy: 0.9982 - val_f1: 0.6108
Epoch 6/7
625/625 [=====] - 4543s 7s/step - loss: 0.0587 - accuracy: 0.9885 - f1: 0.5299 - val_loss: 0.0188 - val_accuracy: 0.9983 - val_f1: 0.6319
Epoch 7/7
625/625 [=====] - 4528s 7s/step - loss: 0.0589 - accuracy: 0.9886 - f1: 0.5231 - val_loss: 0.0176 - val_accuracy: 0.9972 - val_f1: 0.3565
```

Model 3

Visual Results

This model was not trained enough, which explains this results. Model is detecting roots correctly, but it is also not detecting them on lighter spot (plant 2/5).

Best f1 = 0.63



Robotics




Simulation Environment

	Corner	X	Y	Z
0	Top Front Left	-0.26	0.13	0.20
1	Top Front Right	0.18	0.13	0.20
2	Top Back Left	-0.26	-0.26	0.20
3	Top Back Right	0.18	-0.26	0.20
4	Bottom Front Left	-0.26	0.13	0.08
5	Bottom Front Right	0.18	0.13	0.08
6	Bottom Back Left	-0.26	-0.26	0.08
7	Bottom Back Right	0.18	-0.26	0.08

- I identified the maximum and minimum X, Y, Z coordinates
- Sent movement commands in all directions to determine robot's range. During robot's movement I recorded robot's positions.
- Displayed maximum/minimum coordinated in a table

Creating a Gym Environment

- Developed a custom environment to integrate the robot's simulation with RL
- Defined the range of movements
- Defined observation space



Reinforcement Learning

The goal is to train an RL agent to move the pipette to any target position.

Test results:

```
Test completed
Average number of steps to target: 53.7
Success rate (accuracy  $\leq$  0.001 m): 100.0%
numActiveThreads = 0
stopping threads
Thread with taskId 0 with handle 00000000000000EE8 exiting
Thread TERMINATED
finished
numActiveThreads = 0
btShutDownExampleBrowser stopping threads
Thread with taskId 0 with handle 000000000000005F0 exiting
Thread TERMINATED
```

Hyperparameters:

Learning Rate: 0.0003

Batch Size: 64

n_steps: 1000

Gamma: 0.99

Best Result:

Reached pipette goal with accuracy < 1mm
with (average 53.7 steps)

Creating a Controller

Move the pipette tip accurately within the working envelope using PID control.

Performing Gains:

$K_p = 0.8$

$K_i = 0.2$

$K_d = 0.4$

PID controller test results:

```
Target reached within 1.0 mm at step 132
numActiveThreads = 0
stopping threads
Thread with taskId 0 with handle 0000000000000918 exiting
Thread TERMINATED
finished
numActiveThreads = 0
btShutDownExampleBrowser stopping threads
Thread with taskId 0 with handle 000000000000052C exiting
Thread TERMINATED
```

Achieved accuracy < 0.001 m

Assumptions



Photo with only plants

- I designed the pipeline with the assumption that the train/validation/test images provided by NPEC would contain only plants in Petri dish area.

Kaggle image →



Train image →



5 plants on every image

My code was designed with assumption that they are going to be 5 plants in every image. However, in some photos this is not true.

If there are fewer than 5 bounding boxes, their images are replaced with placeholders.

I divide the predicted mask into 5 sections to receive consistent output.

```
for i, (filename, mask) in enumerate(masks):
    print(f"Processing mask {i + 1}/{len(masks)} ({filename})...")

    # generating bounding boxes for detected objects
    bboxes = generate_bounding_boxes(mask, min_area=230, y_range=(370, 730))

    # selecting 5 largest bounding boxes by height
    top_bboxes = get_top_n_bounding_boxes(bboxes, n=5)

    # dividing mask into 5 regions
    regions = divide_mask_into_regions(mask, n=5)

    for region_idx, region in enumerate(regions, start=1):
        # filtering bounding boxes that overlap
        region_bboxes = filter_bboxes_by_region(top_bboxes, region)

        if region_bboxes:
            save_bounding_box_contents(mask, region_bboxes, output_dir=output_path, base_filename=filename, region_number=region_idx)
        else:
            # saving a placeholder if the region has no bounding boxes
            save_path = os.path.join(output_path, f"{filename}_plant_{region_idx}.png")
            with open(save_path, 'w') as f:
                print(f"Empty region {region_idx}, placeholder saved to {save_path}")
```


Limitations

- I don't have GPU on my computer -> Long training times locally. I wasn't able to test multiple models for long enough.
- Big/heavy dataset -> While trying to train my model on Paperspace, I faced issues with uploading data, difficulty transferring data.
- Not trained enough model -> Gaps between roots. Because of it I wasn't able to detect all plants correctly on some images.

Next Steps



- Val_f1 score ≥ 10
- Correctly predicted roots
- Training new model on a server
- Using the entire dataset
- Potentially using the model from notebook 7 and addressing the problem

Developing better model

Creating a PID Controller

- Working on the server provided by BUas
- Designed a PID controller to move pipette tip with precision $\leq 1\text{mm}$
- Implementing three PID controllers responsible for three dimensions (X, Y, Z)

Integrating the Computer Vision Pipeline

- Connecting CV pipeline to PID and RL controllers for precise control
- Converting root tip pixel coordinates into robot-compatible coordinates.
 - Testing the integration

Performance Benchmarking

- Quantifying and comparing performance of PID and RL controllers based on accuracy and speed
- Identifying metrics and methods to evaluate both controllers

Summary



**Thank you for your
attention**

